

NEW YORK UNIVERSITY

MASTER'S THESIS

An L-BFGS-B-NS Optimizer for Non-Smooth Functions

Author:

Wilmer Henao

Supervisor:

Michael L. Overton

*A thesis submitted in fulfillment of the requirements
for the degree of Master of Science in Scientific Computing*

in the

Courant Institute of Mathematical Sciences
Department of Mathematics

May 2014

Declaration of Authorship

I, Wilmer Henao, declare that this thesis titled, 'An L-BFGS-B-NS Optimizer for Non-Smooth Functions' and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

Date:

NEW YORK UNIVERSITY

Abstract

Department of Mathematics

Master of Science in Scientific Computing

An L-BFGS-B-NS Optimizer for Non-Smooth Functions

by Wilmer Henao

This thesis investigates a large scale L-BFGS-B optimizer for smooth functions and how it can be modified to optimize non-smooth functions. The new code is called L-BFGS-B-NS. The changes required include a relaxation of the Wolfe condition, some other changes to the line search algorithm and changes to the termination condition. Some experiments illustrate the results applied to a non-smooth function.

Acknowledgements

I would like to thank my advisor Michael Overton for all the hours of hard work and for all the great recommendations and changes that led to this thesis.

I also would like to thank Allan Kaku and Anders Skajaa for earlier contributions to other versions of the code, and Jorge Nocedal, Ciyu Zhu, Richard Byrd and Peihuang Lu, for letting us change their original code.

Finally, I would like to thank High Performance Computing at NYU for the computing resources and their helpful assistance.

Contents

Declaration of Authorship	i
Abstract	ii
Acknowledgements	iii
Contents	iv
List of Figures	v
List of Tables	vi
1 Introduction	1
2 L-BFGS-B	3
2.1 BFGS	3
2.2 L-BFGS	4
2.3 L-BFGS-B	4
2.3.1 Gradient Projection	4
2.3.2 Subspace Minimization	5
3 Modifications to the L-BFGS-B Algorithm	7
3.1 The Armijo and Wolfe conditions	7
3.2 The Line Search Methodology	8
3.3 The Termination Condition	9
3.3.1 Termination Condition Sub-algorithm	9
3.3.2 The Solution of the Quadratic Program	10
4 Experimental Results	11
4.1 Exit Messages	11
4.2 Modified Rosenbrock Function	12
4.2.1 Performance of L-BFGS-B and L-BFGS-B on Smooth and Non-smooth cases	13
4.2.2 Performance of L-BFGS-B-NS	15

List of Figures

2.1	Graphical Representation of Gradient Projection	5
3.1	Representation of the Armijo Condition in a Nutshell	7
3.2	The Idea behind the Wolfe Condition	8

List of Tables

4.1	Modified Rosenbrock with $p = 2$	14
4.2	Modified Rosenbrock with $p = 1$	15
4.3	Number of algorithm Iterations Changing p	16
4.4	Selected Runs $p = 0.9$	16

Dedicated to my mother and Dr. Ian Malcolm

Chapter 1

Introduction

The problem addressed is to find a local minimizer of the non-smooth minimization problem

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & f(x) \\ \text{s.t.} \quad & l_i \leq x_i \leq u_i, \\ & i = 1, \dots, n. \end{aligned} \tag{1.1}$$

where $f: \mathbb{R}^n \rightarrow \mathbb{R}$, is continuous but not differentiable everywhere and n is large.

The L-BFGS-B algorithm [BLNZ95] is a standard method for solving large instances of (1.1) when f is a smooth function, typically twice differentiable. The name BFGS stands for Broyden, Fletcher, Goldfarb and Shanno, the originators of the BFGS quasi-Newton algorithm for unconstrained optimization discovered and published independently by them in 1970 [Bro70, Fle70, Gol70, Sha70]. This method requires storing and updating a matrix which approximates the inverse of the Hessian $\nabla^2 f(x)$ and hence requires $\mathcal{O}(n^2)$ operations per iteration. The L-BFGS variant [Noc80], where the L stands for “Limited-Memory” and also for “Large” problems, is based on BFGS but requires only $\mathcal{O}(n)$ operations per iteration, and less memory. Instead of storing the $n \times n$ Hessian approximations, L-BFGS stores only m vectors of dimension n , where m is a number much smaller than n . Finally, the last letter B in L-BFGS-B stands for bounds, meaning the lower and upper bounds l_i and u_i in equation (1.1). The L-BFGS-B algorithm is implemented in a FORTRAN software package [ZBNM11].

In this thesis, we first give a brief description of the L-BFGS-B algorithm at a high level and then we introduce a modified algorithm which is more suitable for functions f which may not be differentiable at their local or global optimal points. We call the new

algorithm L-BFGS-B-NS where NS stands for non-smooth. These changes were implemented in a modified version of the FORTRAN code [Hen14] which can be downloaded from a web repository. We report on some numerical experiments that strongly suggest that the new code should be useful for the non-smooth bound-constrained optimization problem (1.1).

We are grateful to Jorge Nocedal and his coauthors for allowing us to modify the L-BFGS-B code and post the modified version.

Chapter 2

L-BFGS-B

This section is a description of the original L-BFGS-B code [ZBNM11, ZBLN97] at a very high level. The original software is intended to find local minimizers of smooth functions. This thesis discusses how to modify the algorithm for non-smooth functions.

2.1 BFGS

BFGS is a standard tool for optimization of smooth functions [NW99]. It is a line search method. The search direction is of type $d = -B_k \nabla f(x_k)$ where B_k is the k^{th} approximation to the inverse Hessian of f .¹ This k^{th} step approximation is calculated via the BFGS formula

$$B_{k+1} = \left(I - \frac{s_k y_k^T}{y_k^T s_k} \right) B_k \left(I - \frac{y_k s_k^T}{y_k^T s_k} \right) + \frac{s_k s_k^T}{y_k^T s_k} \quad (2.1)$$

where $y_k = \nabla f(x_{k+1}) - \nabla f(x_k)$ and $s_k = x_{k+1} - x_k$. BFGS exhibits super-linear convergence on generic problems but it requires $\mathcal{O}(n^2)$ operations per iteration [NW99].

In the case of non-smooth functions, BFGS typically succeeds in finding a local minimizer [LO13]. However, this requires some attention to the line search conditions. These conditions are known as the Armijo and weak Wolfe line search conditions and they are a set of inequalities used for the computation of an appropriate step length that reduces the objective function “sufficiently”. These inequalities will be explained later in section 3.1.

¹When it is exactly the inverse Hessian the method is known as Newton’s method. Newton’s method has quadratic convergence but requires the explicit calculation of the Hessian at every step.

2.2 L-BFGS

L-BFGS stands for Limited-memory BFGS. This algorithm approximates BFGS using only a limited amount of computer memory to update an approximation to the inverse of the Hessian of f . Instead of storing a dense $n \times n$ matrix, L-BFGS keeps a record of the last m iterations where m is a small number that is chosen in advance². For this reason the first m iterations of BFGS and L-BFGS produce exactly the same search directions if the initial approximation B_0 is set to the identity matrix.

Because of this construction, the L-BFGS algorithm is less computationally intensive and requires only $\mathcal{O}(mn)$ operations per iteration. So it is much better suited for problems where the number of dimensions n is large.

2.3 L-BFGS-B

Finally L-BFGS-B is an extension of L-BFGS. The B stands for the inclusion of Boundaries. L-BFGS-B requires two extra steps on top of L-BFGS. First, there is a step called *gradient projection* that reduces the dimensionality of the problem. Depending on the problem, the gradient projection could potentially save a lot of iterations by eliminating those variables that are on their bounds at the optimum reducing the initial dimensionality of the problem and the number of iterations and running time. After this *gradient projection* comes the second step of *subspace minimization*. During the *subspace minimization* phase, an approximate quadratic model of (1.1) is solved iteratively in a similar way that the original L-BFGS algorithm is solved. The only difference is that the step length is restricted as much as necessary in order to remain within the *lu*-box defined by equation (1.1).

2.3.1 Gradient Projection

The L-BFGS-B algorithm was designed for the case when n is large and f is smooth. Its first step is the gradient projection similar to the one outlined in [CGT88, MT89] which is used to determine an active set corresponding to those variables that are on either their lower or upper bounds. The active set defined at point x^* is:

$$\mathcal{A}(x^*) = \{i \in \{1 \dots n\} | x_i^* = l_i \vee x_i^* = u_i\} \quad (2.2)$$

²In this thesis $m < 20$, and in practice numbers between 5 and 10 are regularly used. There is no way of knowing *a priori* what choice of m will provide the best results as will be illustrated later.

Working with this active set is more efficient in large scale problems. A pure line search algorithm would have to choose a step length short enough to remain within the box defined by l_i and u_i . So if at the optimum, a large number \mathcal{B} of variables are either on the lower or the upper bound, as many as \mathcal{B} of iterations might be needed. Gradient projection tries to reduce this number of iterations. In the best case, only 1 iteration is needed instead of \mathcal{B} .

Gradient projection works on the linear part of the approximation model:

$$m_k(x) = f(x_k) + \nabla f(x_k)^T(x - x_k) + \frac{(x - x_k)^T H_k(x - x_k)}{2} \quad (2.3)$$

where H_k is a L-BFGS-B approximation to the Hessian $\nabla^2 f$ stored in the implicit way defined by L-BFGS.

In this first stage of the algorithm a piece-wise linear path starts at the current point x_k in the direction $-\nabla f(x_k)$. Whenever this direction encounters one of the constraints, the path turns corners in order to remain feasible. The path is nothing but the feasible piece-wise projection of the negative gradient direction on the constraint box determined by the values l and u . At the end of this stage, the value of x that minimizes $m_k(x)$ restricted to this piece-wise gradient projection path is known as the “Cauchy point” x^c . See Figure 2.1.

2.3.2 Subspace Minimization

The problem with gradient projection is that its search direction does not take advantage of information provided implicitly by the Hessian H_k , and therefore the speed of convergence is at best linear. It is for this reason that a second stage is necessary. Stage 2 (subspace minimization) uses an L-BFGS implicit approximation of the inverse Hessian matrix restricted to the free variables that are not in the active set $\mathcal{A}(x^c)$.

The starting position for stage 2 will be the previously found Cauchy point and the goal is to find a new $\bar{x} = x^c + \alpha^* \hat{d}$. The idea at a higher level is to minimize (2.3) over the free variables subject to their lower and upper bounds. First, the L-BFGS algorithm provides a new search direction \hat{d}^u of the *unconstrained* problem that takes implicit advantage of approximations of the Hessian matrix restricted to the free variables. After an unconstrained search direction has been found, the constraints are taken into account and the search direction is restricted to the l, u bounding box via a step length factor α^* . The step length is chosen so that the new point \bar{x} satisfies the Armijo and Wolfe³

³The Armijo and Wolfe conditions will be explained on section 3.1




FIGURE 2.1: The arrow represents the direction of the negative gradient. The dotted path represents the projected gradient path. The contours represent the level sets of the model. The optimal point (the '*' in red) is the Cauchy point x^c

conditions. A restriction on the step length is added so that the next iteration stays feasible. Sometimes it is not possible to satisfy the Wolfe condition due to the bounded nature of the problem, so in these cases, only the Armijo condition needs to be satisfied. Once this step length is found, the next step is to check the termination condition. If the termination condition fails, a new gradient projection and subspace minimization will be needed and the method repeats. If the termination condition is successful, the program exits with an appropriate exit message.

Chapter 3

Modifications to the L-BFGS-B Algorithm

We made three main changes to the original L-BFGS-B algorithm. They concern the line search Wolfe conditions, the line search methodology, and the termination condition.

3.1 The Armijo and Wolfe conditions

It is accepted that the Armijo and Wolfe conditions work very well whenever the function f is smooth [LF01]. The Armijo condition, also known as the sufficient decrease requirement in the direction d_k , is defined as

$$f(x_k + \alpha_k d_k) \leq f(x_k) + c_1 \alpha_k d_k^T \nabla f(x_k) \quad (3.1)$$

where $0 < c_1 < 1$ is a constant, often $c_1 = 10^{-4}$ [NW99]. This condition guarantees “sufficient decrease” of the function. It is possible to continue decreasing without ever reaching the optimum if the Armijo condition is not required as is shown in Figure 3.1.

The other condition, which is the one that was actually changed, is the curvature condition, of which the most popular version is the “strong Wolfe” curvature condition:

$$|d_k^T \nabla f(x_k + \alpha_k d_k)| \leq c_2 |d_k^T \nabla f(x_k)| \quad (3.2)$$

Here d_k represents the search direction and c_2 is a constant such that $0 < c_1 < c_2 < 1$; often $c_2 = 0.9$ [NW99]. The strong Wolfe condition is a natural choice for optimization




FIGURE 3.1: In this figure, the iterations always reduce the value of the function a little bit, but never enough to go below 12




FIGURE 3.2: The logic of the Wolfe conditions is this. Starting at point A, Point B is a step in the right direction, however, point C offers a “flatter” tangent and should be closer to the optimum which has a tangent of zero (Smooth case).

of smooth functions. Its goal is to find a step length long enough that the slope has been reduced “sufficiently” as illustrated in figure 3.2, but the problem is that the condition, as it is, does not work well for the non-smooth case. This is because near the minimal points there may be abrupt changes in the gradient. A good example of this problem is the function $f(x) = |x|$, where the slope never becomes flat near the optimal point.

The weak Wolfe condition defined as

$$d_k^T \nabla f(x_k + \alpha_k d_k) \geq c_2 d_k^T \nabla f(x_k) \quad (3.3)$$

can be used in the non-smooth case. It is all that is needed to guarantee that the BFGS updated inverse Hessian approximation is positive definite [NW99]. This weak version is suited for the problems in this thesis and it was implemented as part of the line search algorithm explained in the next section.

3.2 The Line Search Methodology

The original FORTRAN software [ZBNM11] contains a line search subroutine. It was partially changed for the purpose of this thesis. The old version of the code was commented out.

The change in the Wolfe condition has already been described, but there is an additional problem with the function *dcstep* in the non-smooth case. The function *dcstep* was designed to work only with smooth functions in mind. The algorithm in *dcstep* takes advantage of quadratic and cubic approximations to the function in order to calculate step lengths that satisfy Armijo and Wolfe conditions. These second and third order approximations do not work well in the non-smooth case, and the optimizer breaks down using the line search as it is.

The solution to this particular issue is to use a line search similar to the one suggested in [LO13] and in [OS12]. This approach is to double the step length while the Armijo condition is violated, and once the interval has been bracketed, do bisection until both the Armijo and Wolfe conditions are satisfied. The only difference with the approach in this thesis is that the line search in HANSO can double its step length up to 30 times, whereas in this thesis, the step length can double only as long as it is less than the maximum value that guarantees feasibility of the solution (the maximum established in the first step of the original line search).

3.3 The Termination Condition

In the case of smooth functions, L-BFGS-B checks whether the algorithm has converged by means of the *projected gradient* which is nothing but the projection of the negative gradient onto the bounding box defined by l and u . If this projected gradient has a small norm the algorithm terminates. In the case of non-smooth functions however, the function at the minimum may have a “wedge”. In this wedge the projected gradient

may not vanish (it is not defined at the “bottom” of the wedge, such as is the case for $f(x) = |x|$ at $x = 0$). Furthermore, if there is a sequence of points that approaches the optimum x in a direction \vec{p} , the projected gradients corresponding to this sequence of points might be completely different from the projected gradients associated with a sequence of points that approach the optimum x from a different direction.

3.3.1 Termination Condition Sub-algorithm

Lewis and Overton formulate an algorithm that gives a practical solution to this problem in section 6.3 of [LO13] in the case of unconstrained non-smooth optimization. They suggest computing the norm of the smallest vector that is part of the convex hull of gradients evaluated at points near the optimum candidate x and terminate if this is sufficiently small. The neighborhood is defined as those points at which the gradient has already been evaluated with a distance to x smaller than a small tolerance $\tau_x > 0$ and no more than $J \in \mathbb{N}$ iterations back in history. This list of gradients is referred to as the set \mathcal{G} [LO13].

With this list \mathcal{G} of gradients at hand, the next step is to find the vector with the minimal norm contained in the convex hull of these gradients. If the minimum such norm is smaller than another tolerance τ_d , the algorithm terminates.

In order to find this vector, there is the need to solve a quadratic problem. Every vector in the convex hull can be expressed as a convex combination Gz of those vectors in \mathcal{G} , where G is the matrix with columns made up of gradients in \mathcal{G} and z is such that $\sum z_i = 1$ and $z_i \geq 0$.

The objective is to find the right combination of z that minimizes the norm $\|Gz\|_2$. This is equivalent to solving the following optimization problem

$$\begin{aligned} \min \quad & q(z) = \|Gz\|_2^2 = z^T G^T G z \\ \text{s.t.} \quad & \sum z_i = 1 \\ & z_i \geq 0. \end{aligned} \tag{3.4}$$

The solution to this problem z^* defines the associated vector Gz^* , so if $\|Gz^*\|_2 < \tau_d$ the algorithm terminates.

In the bound-constrained case, it is important to notice that instead of the gradient we have to work with the projected gradient. In the unconstrained case, if a component of the gradient is not zero this yields a direction of descent. But in the bounded case, it may be impossible to reduce f in this direction because the boundary may have been

reached. For this reason the gradients have to be projected onto the bounding box, and it is these projected gradients that we incorporate in the termination condition of L-BFGS-B-NS.

3.3.2 The Solution of the Quadratic Program

The solution of the quadratic program (3.4) is obtained using a practical primal-dual method. This is the same method implemented by Skajaa [Ska10] in his thesis. His code `qpspecial` was implemented in FORTRAN for this thesis. The method is the well known Mehrotra's Predictor-Corrector algorithm applied to quadratic programming, as explained in chapter 16 of [NW99].

Chapter 4

Experimental Results

The L-BFGS-B implementation was tested on the high performance cluster machines at NYU. In order to run these tests it was necessary to create a series of PBS files¹ using a PBS generator script. This script generator created PBS files which in turn run bash shell scripts². Several of these shell scripts are available at the repository [Hen14]. The main reason to run scripts this way is because it achieves parallelism, and because the system sends confirmation e-mails and statistics about the different stages of the processes giving a lot of control to the practitioner.

4.1 Exit Messages

The original L-BFGS-B optimizer displays different messages depending on the condition that triggered the exit. The following is a list of some of the most common exit messages in the original L-BFGS-B optimizer.

- “ABNORMAL_TERMINATION_IN_LNSRCH” This message means that there was a problem and the program’s exit was premature. It is typically found for non-smooth functions where the line search breaks down. But the message could also be symptomatic of other problems.
- “CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_LT_PGTOL”: Means that convergence was achieved because the norm of the projected gradient is small

¹PBS stands for Portable Batch System. This is software that performs job scheduling. It is used by High Performance Computing at NYU (and many other High Performance Computing Centers) to allocate computational tasks. In order to run jobs at the high performance clusters, a series of PBS batch files need to be created.

²Bash is a command processor. Each Bash script that was created includes a series of computer commands, namely, execution of the original L-BFGS-B software and the new code L-BFGS-B-NS.

enough. Notice that this convergence message does not apply to L-BFGS-B-NS because of particular requirements for non-smooth functions involving the convex hull of projected gradients as explained in section 3.3. Instead it is replaced by

- “CONVERGENCE: ZERO_GRAD_IN_CONV_HULL” This means that the termination condition discussed in section 3.3 was satisfied³.
- “CONVERGENCE: REL_REDUCTION_OF_F_LT_FACTR*EPSMCH”: This convergence condition is achieved whenever the relative reduction of the value of function f is smaller than a predefined factor times the machine precision ϵ . This exit message does not apply to our tests. It was disabled by setting the factor “FACTR” to zero, both in our runs of L-BFGS-B-NS and our tests using the original code L-BFGS-B.

The limit on the number of iterations was set to 10,000.

4.2 Modified Rosenbrock Function

Consider a modified version of the Rosenbrock function problem [Ros61]:

$$f(x) = (x_1 - 1)^2 + \sum_{i=2}^n |x_i - x_{i-1}^2|^p \quad (4.1)$$

We can study the properties of function f based on the properties of the function $\phi(t_i)$, where $\phi(t_i) = |t_i|^p$ and $t_i = x_i - x_{i-1}^2$. The properties of the function depend on the value of the p parameter⁴. This function can be proven to be locally Lipschitz continuous whenever $p \geq 1$. However, its second derivative blows up at zero whenever $p < 2$. Note that although $\phi(t_i)$ is convex for $p \geq 1$, f is not convex.

The properties of $\phi(t_i)$ can be separated into different cases. Whenever $p > 1$ the derivative can be represented as:

$$\frac{d}{dt}\phi(t) = \pm p|t|^{p-1} \quad (4.2)$$

and therefore, the limit of the derivative exists and is equal to zero near $t = 0$:

$$\lim_{t \rightarrow 0} \frac{d}{dt}\phi(t) = 0$$

³This does not mean that the resulting vector is exactly equal to zero 0, but it is small enough to satisfy the termination condition.

⁴The original Rosenbrock function had a value of $p = 2$ and the second term was multiplied by 100.

From this we conclude that f has a smooth first derivative for $p > 1$.

However, if $p = 1$, $\phi(t) = |t|$, and the absolute value function is not differentiable at $t = 0$. Note that, $\phi(t)$ is Lipschitz continuous at $t = 0$.

The second derivative provides a bit more of information.

$$\frac{d^2}{dt^2}\phi(t) = \pm p(p-1)|t|^{p-2} \quad (4.3)$$

If $p \geq 2$ the function is twice continuously differentiable. However if $p < 2$, the second derivative becomes $\frac{p(p-1)}{|t|^q}$, where $q = 2 - p > 0$, and this second derivative blows up as $|t| \rightarrow 0$. The special case $p = 1$ has second derivative equal to zero since $p(p-1) = 0$ except at $t = 0$ where it is undefined. For $p < 1$, ϕ is not Lipschitz continuous at $t = 0$.

Having explained the characteristics of the function, the next thing that needs to be defined is the region to be tested. We chose the region to be defined by the “box” with boundaries

$$x_i = \begin{cases} [-100, 100] & \text{if } i \in \text{even numbers} \\ [10, 100] & \text{if } i \in \text{odd numbers} \end{cases} \quad (4.4)$$

The initial point was chosen to be the midpoint of the box, plus a different small perturbation for each dimension, chosen so that the line search does not reach the boundary of several dimensions in one step:

$$x_i = \frac{u_i + l_i}{2} - (1 - 2^{1-i}) \quad (4.5)$$

The problem is twice continuously differentiable for values of $p \geq 2$, but as the values of p approach 1, the original L-BFGS-B optimizer should start to have problems. We tested the original L-BFGS-B optimizer on the modified Rosenbrock function with p varying between 2 and 0.9.

For a value of $p = 2$, the original L-BFGS-B yields good results as seen in Table 4.1.

4.2.1 Performance of L-BFGS-B and L-BFGS-B on Smooth and Non-smooth cases

In order to compare the results of L-BFGS-B with the results of L-BFGS-B-NS, we changed the norm of the gradient in L-BFGS-B from the infinity to the euclidean norm, other

m	n	L-BFGS-B results				L-BFGS-B-NS results			
		Iter.	#fg	f	NPG	Iter.	#fg	f	NSVCHPG
5	4	2	23	9305.93	0.00207341	10	16	9305.93	3.57E-005
10	4	2	29	9305.93	0.0034576345	10	16	9305.93	6.07E-005
20	4	2	30	9305.93	0.0011789284	10	16	9305.93	6.07E-005
5	6	2	19	18531.14	0.0045643053	10	16	18531.14	4.92E-004
10	6	2	19	18531.14	0.0098673465	10	16	18531.14	1.89E-004
20	6	2	19	18531.14	0.0098672632	10	16	18531.14	1.89E-004
5	8	2	21	27756.35	0.0046836051	13	19	27756.35	4.02E-004
10	8	2	20	27756.35	0.0094596482	14	20	27756.35	2.38E-006
20	8	2	20	27756.35	0.0094708831	14	20	27756.35	2.61E-006
5	10	2	21	36981.56	0.0033234587	15	21	36981.56	3.92E-005
10	10	2	21	36981.56	0.0081012695	15	21	36981.56	8.65E-005
20	10	2	21	36981.56	0.0082688034	15	21	36981.56	1.42E-004
5	20	2	20	83107.61	0.0027406376	16	21	83107.61	1.41E-005
10	20	2	20	83107.61	0.0041252207	16	21	83107.61	3.36E-005
20	20	2	20	83107.61	0.0041233412	16	21	83107.61	7.15E-005
5	100	2	20	452116.01	0.0019111682	16	21	452116.01	1.45E-004
10	100	2	20	452116.01	0.002587029	16	21	452116.01	3.29E-004
20	100	2	20	452116.01	0.0025871555	16	21	452116.01	6.10E-004
5	200	2	20	913376.52	0.0023793373	17	22	913376.52	1.64E-004
10	200	2	20	913376.52	0.0031088411	17	22	913376.52	3.51E-004
20	200	2	20	913376.52	0.003108401	17	22	913376.52	6.81E-004
5	1000	2	20	4603460.52	0.0035296465	18	24	4603460.52	4.04E-005
10	1000	2	20	4603460.52	0.0045067318	17	22	4603460.52	3.24E-004
20	1000	2	20	4603460.52	0.0045070181	18	30	4603460.52	1.18E-005

TABLE 4.1: Satisfactory results for the original algorithm L-BFGS-B applied to the Modified Rosenbrock function with $p = 2$. And results for L-BFGS-B-NS; NPG: Norm of projected Gradient with tolerance 10^{-6} . NSVCHPG: Norm of Smallest Vector in Convex Hull of Projected Gradients with $\tau_d = 10^{-6}, \tau_x = 10^{-3}$

than that L-BFGS-B is exactly the same version as the original. In order to keep things comparable, the tolerance for the norm of the projected gradient in L-BFGS-B and τ_d from L-BFGS-B-NS were both set to 10^{-6} .

The values of m (the memory of L-BFGS) that were tested are 5, 10 and 20. The number of dimensions in this exercise ranges from 6 to 1000. The column $\#fg$ stands for the number of function and gradient evaluations taken, f stands for the optimal value that was achieved by the optimization. NPG shows the norm of the projected gradient, the termination tolerance for the euclidean norm of the projected gradients was 10^{-6} . In all cases this test was satisfied.

In table 4.1 $NSVCHPG$ is the norm of the smallest vector in the convex hull of projected gradients. You can see that since this function is smooth when $p = 2$, L-BFGS-B has no problems solving it and that L-BFGS-B-NS reaches exactly the same value.

The overall conclusion from this exercise is that the original L-BFGS-B optimizer works well for the smooth modified Rosenbrock case and that L-BFGS-B-NS has a performance that can achieve a minimum that is just as good.

On the other hand, the value of $p = 1$ leads to an abnormal line search termination for L-BFGS-B in all of the cases presented. This is to be expected as the function is non-smooth. See table 4.2 where the norm of the resulting projected gradient never

m	n	p	L-BFGS-B results			L-BFGS-B-NS results		
			#fg	f	NPG	#fg	f	NSVCHPG
10	4	1	68	274.68	96.84	93	177.00	4.77E-09
20	4	1	68	274.68	96.84	93	177.00	4.77E-09
10	6	1	57	371.80	96.81	113	274.68	6.54E-08
20	6	1	57	371.80	96.81	113	274.68	6.54E-08
10	8	1	59	468.38	96.84	106	371.52	2.42E-09
20	8	1	59	468.38	96.84	106	371.52	2.42E-09
10	10	1	59	565.28	96.83	188	521.64	1.98E-08
20	10	1	59	565.28	96.83	188	521.64	1.98E-08
10	20	1	69	1,049.37	96.84	188	952.54	1.89E-08
20	20	1	69	1,049.37	96.84	279	952.54	7.90E-08
10	100	1	55	4,923.83	96.83	138	4,826.05	5.56E-08
20	100	1	55	4,923.83	96.83	138	4,826.05	5.56E-08
10	200	1	55	9,767.31	96.83	140	9,667.93	7.72E-08
20	200	1	55	9,767.31	96.83	112	9,667.93	5.18E-08
10	1000	1	55	48,515.21	96.83	163	48,403.02	1.70E-07
20	1000	1	55	48,515.21	96.83	177	48,403.02	1.80E-07

TABLE 4.2: Unsatisfactory results for the original algorithm L-BFGS-B applied to the Modified Rosenbrock function with $p = 1$. And converging results for L-BFGS-B-NS; NPG: Norm of projected Gradient with tolerance = 10^{-6} . NSVCHPG: Norm of Smallest Vector in Convex Hull of Projected Gradients with $\tau_d = 10^{-6}$, $\tau_x = 10^{-3}$

approaches zero. In this exercise, the memory length m of L-BFGS, does not have an impact on the final value f of the optimization, but this is because all cases crashed before the 5th iteration and therefore all different cases of m end up looking exactly the same.

4.2.2 Performance of L-BFGS-B-NS

Several other values of p were also tested, among others 1.5, 1.1, 1.01, 1.001, ... , 1.00001, 1. With this tolerance L-BFGS-B always crashes as expected, those values where p is closer to 1 are the most difficult for the original algorithm to handle. Values generated via L-BFGS-B-NS are comparatively better whenever $p < 2$, since the function is “less” smooth. Some runs of L-BFGS-B-NS fail to converge using the termination condition from section 3.3.

In table 4.3, the parameter p is varied and all other parameters are held constant.

And finally a test run for values $p < 1$. In particular here a test of $p = 0.9$ on Table 4.4. This test fails because of Lipschitz discontinuity of the function in the region.

p	Iterations	Value of f	Norm Smallest Vector in Convex Hull
2	8	41,116,905.61	1.88E-03
1.1	24	764,853.32	1.72E-06
1.0001	27	484,394.49	1.91E-08
1.00001	84	484,195.01	1.06E-06
1.0000001	21	484,173.43	1.77E-08

TABLE 4.3: This is the number of algorithm iterations for different values of p . The value of the projected gradient is presented as well. This exercise was run with $n = 10,000$, $m = 10$ and $\tau_d = \tau_x = 10^{-3}$

m	n	p	L-BFGS-B results			L-BFGS-B-NS results	
			#fg	f	NFG	#fg	f
5	2	0.9	21	142.3895656804	97	50358	81
10	2	0.9	21	142.3895656804	97	50358	81
20	2	0.9	21	142.3895656804	97	50358	81
5	4	0.9	57	203.6867061962	96.8244771931	20010	203.6717690175
10	4	0.9	57	203.6867061962	96.8244771931	20007	201.3635208544
20	4	0.9	57	203.6867061962	96.8244771931	20007	201.3635208544
5	6	0.9	90	264.8904504715	96.8377187807	20012	264.9727640136
10	6	0.9	90	264.8904504715	96.8377187807	20006	263.4634677358
20	6	0.9	90	264.8904504715	96.8377187807	20006	263.2710714257
5	8	0.9	65	326.2465886362	96.8354703277	588470	326.2699902292
10	8	0.9	65	326.2465886362	96.8354703277	20003	324.2733728009
20	8	0.9	65	326.2465886362	96.8354703277		
5	10	0.9	55	387.6474214791	96.8328164056	588173	387.5671220185
10	10	0.9	55	387.6474214791	96.8328164056	19998	386.9547954789
20	10	0.9	55	387.6474214791	96.8328164056	20002	386.5111464469
5	20	0.9	55	695.4247337795	96.8148669898	588234	694.0527666868
10	20	0.9	55	695.4247337795	96.8148669898	20005	694.0391973604
20	20	0.9	55	695.4247337795	96.8148669898	20003	682.9908544633
5	50	0.9	53	1631.27559033	96.7138084598	588409	1613.5096637579
10	50	0.9	53	1631.27559033	96.7138084598	20006	1613.5043396518
20	50	0.9	53	1631.27559033	96.7138084598	20002	1604.0578060639
5	100	0.9	53	3183.6955082727	96.7088482124	588409	3145.9378051899
10	100	0.9	53	3183.6955082727	96.7088482124	20005	3145.9332306031
20	100	0.9	53	3183.6955082727	96.7088482124	20007	3144.6121183855
5	200	0.9	53	6288.5762389464	96.7063705729	588409	6210.7940850592
10	200	0.9	53	6288.5762389464	96.7063705729	20010	6210.7940838524
20	200	0.9	53	6288.5762389464	96.7063705729	20007	6209.6424888061
5	1000	0.9	53	31127.7188550318	96.7043901178	588764	30729.6443168679
10	1000	0.9	53	31127.7188550318	96.7043901178	10735312	30729.6443166712
20	1000	0.9	53	31127.7188550318	96.7043901178	20013	30729.6408804572

TABLE 4.4: Collection of selected runs for both L-BFGS-B and L-BFGS-B-NS. None of them show convergence since this function is not Lipschitz continuous.

Bibliography

- [BLNZ95] Richard H. Byrd, Peihuang Lu, Jorge Nocedal, and Ci You Zhu. A limited memory algorithm for bound constrained optimization. *SIAM J. Sci. Comput.*, 16(5):1190–1208, 1995.
- [Bro70] C. G. Broyden. The convergence of a class of double-rank minimization algorithms. II. The new algorithm. *J. Inst. Math. Appl.*, 6:222–231, 1970.
- [CGT88] A. R. Conn, N. I. M. Gould, and Ph. L. Toint. Global convergence of a class of trust region algorithms for optimization with simple bounds. *SIAM J. Numer. Anal.*, 25(2):433–460, 1988.
- [Fle70] R. Fletcher. A new approach to variable metric algorithms. *The Computer Journal*, 13(3):317–322, 1970.
- [Gol70] Donald Goldfarb. A family of variable-metric methods derived by variational means. *Math. Comp.*, 24:23–26, 1970.
- [Hen14] Wilmer Henao. L-BFGS-B-NS. <https://github.com/wilmerhenao/L-BFGS-B-NS>, 2014.
- [LF01] Dong-Hui Li and Masao Fukushima. On the global convergence of the BFGS method for nonconvex unconstrained optimization problems. *SIAM J. Optim.*, 11(4):1054–1064 (electronic), 2001.
- [LO13] Adrian S. Lewis and Michael L. Overton. Nonsmooth optimization via quasi-Newton methods. *Math. Program.*, 141(1-2, Ser. A):135–163, 2013.
- [MT89] Jorge J. Moré and Gerardo Toraldo. Algorithms for bound constrained quadratic programming problems. *Numer. Math.*, 55(4):377–400, 1989.
- [Noc80] Jorge Nocedal. Updating quasi-Newton matrices with limited storage. *Math. Comp.*, 35(151):773–782, 1980.
- [NW99] Jorge Nocedal and Stephen J. Wright. *Numerical optimization*. Springer Series in Operations Research. Springer-Verlag, New York, 1999.

- [OS12] Michael Overton and Anders Skajaa. Hanso 2.02. <http://www.cs.nyu.edu/faculty/overton/software/hanso/>, 2012.
- [Ros61] H. H. Rosenbrock. An automatic method for finding the greatest or least value of a function. *Comput. J.*, 3:175–184, 1960/1961.
- [Sha70] D. F. Shanno. Conditioning of quasi-Newton methods for function minimization. *Math. Comp.*, 24:647–656, 1970.
- [Ska10] Anders Skajaa. Limited memory BFGS for nonsmooth optimization. Master’s thesis, New York University, Courant Institute of Mathematical Sciences, 251 Mercer Street, New York, NY 10012, 2010.
- [ZBLN97] Ciyou Zhu, Richard H. Byrd, Peihuang Lu, and Jorge Nocedal. Algorithm 778: L-BFGS-B: Fortran subroutines for large-scale bound-constrained optimization. *ACM Trans. Math. Software*, 23(4):550–560, 1997.
- [ZBNM11] Ciyou Zhu, Richard Byrd, Jorge Nocedal, and Jose Luis Morales. L-BFGS-B 3.0. <http://www.ece.northwestern.edu/~nocedal/lbfgsb.html>, 2011.